

Computational Semantics



CS224N

Christopher Manning

(Borrows some slides from Mary Dalrymple,
Jason Eisner, and Jim Martin)



Why study computational semantics?

- Because everyone has been wanting to hear about this all course!?
- Obvious high-level applications
 - Summarization
 - Translation
 - Question answering
 - Information access
 - Talking to your pet robot
 - Speech user interfaces
- The next generation of intelligent applications need deeper semantics than we have seen so far
 - Often you must understand well to be able to act




Shallow vs. deep semantics

- We can do more than one might have thought without deep linguistic analysis
 - This is *the* lesson of the last decade
- But we can't do everything we would like:
 - Not all tasks can ignore higher structure
 - Unsuitable if new text must be generated
 - Unsuitable if machine must act rather than relying on user to interpret material written by the author of the document
- You get what you pay for:
 - Cheap, fast, low-level techniques are appropriate in domains where speed and volume are more important than accuracy
 - More computationally expensive, higher-level techniques are appropriate when high-quality results are required



SHRDLU

<http://hci.stanford.edu/winograd/shrdlu/>





Google: What is the capital of Algeria?

what is the capital of algeria - Google Search

www.google.com/#hl=en&safe=off&tbo=d&client=psy-ab&q=what+is+the+capital+of+algeria&oq=What+is+the+capital+of+algeria

Search Images Drive Sites Groups Contacts Maps YouTube News More

Google what is the capital of algeria manning@stanford.edu

Web Images Maps Shopping Videos More Search tools

About 46,900,000 results (0.30 seconds)

Algiers

Algeria, Capital

[Algeria - Wikipedia, the free encyclopedia](#)
en.wikipedia.org/wiki/Algeria
Its **capital** (and most populous city) is **Algiers**. The territory of today's **Algeria** was the home of many ancient cultures, including Aterian and Capsian cultures.

[Algiers - Wikipedia](#)
en.wikipedia.org/wiki/Algiers
Algiers (Arabic: الجزائر, al-Jazā'ir; **Algerian** Arabic: الجزائر; Berber: Dzayer / ⵟⵖⵣⵉⵔ; German: Alger) is the **capital** and largest city of **Algeria**. According to the ...

[What is the capital city of Algeria](#)
wiki.answers.com > ... > Countries States and Cities > Algeria
What is the capital city of Algeria? In: **Algeria** [Edit categories]. Answer: == Answer == The **capital of Algeria** is **Algiers**. Improve answer. Contributor: Kadia ...

[What is the capital of Algeria - Wiki Answers](#)
wiki.answers.com > ... > Countries States and Cities > Algeria
What is the capital of Algeria? In: **Algeria** [Edit categories]. Answer: Alger **Algiers**
The **capital** is **Algiers** , and also the largest city in **Algeria** . Improve answer ...

Algiers

Algiers is the capital and largest city of Algeria. According to the 1998 census, the population of the city proper was 1,519,570 and that of the urban agglomeration was 2,135,630. In 2009, the population was about 3,500,000. Wikipedia

Weather: 66°F (19°C), Wind S at 5 mph (8 km/h), 52% Humidity

Local time: Monday 6:11 PM

Points of interest: Notre Dame d'Afrique, Martyrs Memorial, More

Feedback / More info

Display a menu



Google: Which countries does the Danube flow through?

A screenshot of a Google search result page. The browser window shows the URL "www.google.com/#hl=en&sugexp=les%3B&gs_nf=3&gs_mss=Which%20countries%20does%20the%20d&qe=V2hpY2ggY291br". The search bar contains the query "Which countries does the danube flow through". Below the search bar, a dropdown menu shows suggestions: "which countries does the danube flow through", "which countries does the danube river flow through", "which countries does the danube river run through", and "which countries does the danube run through". The search results are displayed below, with the first result being "Danube - Wikipedia, the free encyclopedia". The second result is "What countries does the Danube River run through" from wiki.answers.com. The third result is "What country does the river Danube flow through" also from wiki.answers.com. The fourth result is "Danube River (river, Europe) -- Britannica Online Encyclopedia". The fifth result is "How many countries does the danube flow through? - Yahoo! UK ...". The page also shows a "SafeSearch off" button and a user profile "manning@stanford.edu".



Google: What are the capitals of the countries bordering the Baltic Sea?

what are the capitals of the countries that border the baltic sea? - Google Search

www.google.com/#hl=en&safe=off&tbo=d&client=psy-ab&q=What+are+the+capitals+of+the+countries+that+border+the

What's r... www.jai... LingPip... named... Why Ini... Optimal... Scaling... The SIL... what ar... >> +

Search Images Drive Sites Groups Contacts Maps YouTube News More -

Google What are the capitals of the countries that border the Baltic Sea? manning@stanford.edu

Web Images Maps Shopping More Search tools SafeSearch off

About 7,790,000 results (0.18 seconds)

[What European countries have borders on the baltic sea](#)
wiki.answers.com > ... > Science > Geography > Borders and Boundaries
9 countries border the Baltic sea - Denmark, Estonia, Finland, Germany, Latvia, Lithuania, Poland, Russia and Sweden. Estonia, Latvia and Lithuania are ...

[What Countries Border the Baltic Sea? | eHow.co.uk](#)
www.ehow.co.uk > Hobbies
Only a small portion of southern Norway **borders the Baltic Sea**. The **country** has an area of 125,097 square miles and a population of 4.4 million. The **capital** city ...


[What are three Scandinavian countries and at least two other ...](#)
www.chacha.com > Categories > Travel
What are three Scandinavian **countries** and at least two other **nations** that **border the baltic sea**? ChaCha Answer: Sweden, Norway, Finla...

[What 4 countries border the Baltic sea on the east side that was that ...](#)
www.chacha.com > Categories > Travel
What 4 **countries** **border the Baltic sea** on the east side that was that was once part of Russia? ChaCha Answer: **Countries** ... Do You Know The State **Capitals**?

Display a menu

What we say to dogs

Okay, Ginger! I've had it!
You stay out of the garbage!
Understand, Ginger? Stay out
of the garbage, or else!



What they hear


Lemon

blah blah GINGER blah
blah blah blah blah
blah blah GINGER blah
blah blah blah blah...






Apple Knowledge Navigator video (1987)



Precise semantics. An early example: Chat-80

- Developed between 1979 and 1982 by Fernando Pereira and David Warren; became Pereira's dissertation
- Proof-of-concept natural language interface to database system
- Used in projects: e.g. Shoptalk (Cohen et al. 1989), a natural language and graphical interface for decision support in manufacturing
- Used in an AppliedNLP-2000 conference paper! [\[Asking about train routes and schedules – still live\]](#)
- Available in /afs/ir/class/cs224n/src/chat
 - UNFORTUNATELY STILL A BIT BUGGY AFTER PROLOG UPGRADES, AND ONLY RUNS ON OLD MYTHS ☹



Things you could ask...

- What is the total area of countries south of the Equator and not in Australasia?
- What is the average area of the countries in each continent?
- Is there more than one country in each continent?
- What are the countries from which a river flows into the Black_Sea?
- Which country bordering the Mediterranean borders a country that is bordered by a country whose population exceeds the population of India?




Chat-80 trace (small)

Question: What is the capital of
Australia?
Parse: 0.0sec.
whq
\$VAR
1
s
np
3+sin
wh(B)
[]
verb(be,active,pres+fin,[],pos)
arg
dir
np
3+sin

np_head
det(the(sin))
[]
capital
pp
prep(of)
np
3+sin
name(australia)
[]
[]

Semantics: 0.0sec.
answer([B]) :-
capital(australia,B)

canberra.



The CHAT-80 Database

% Facts about countries.

% country(Country,Region,Latitude,Longitude,

% Area (sqmiles), Population, Capital,Currency)

country(andorra,southern_europe,42,-1,179,
25000,andorra_la_villa,franc_peseta).

country(angola,southern_africa,-12,-18,481351,
5810000,luanda,?).

country(argentina,south_america,-35,66, 1072067,
23920000,buenos_aires,peso).

capital(C,Cap) :- country(C,_,_,_,_,_,Cap,_).



The CHAT-80 grammar

```
/* Sentences */
```

```
sentence(S) --> declarative(S), terminator(.).
```

```
sentence(S) --> wh_question(S), terminator(?).
```

```
sentence(S) --> yn_question(S), terminator(?).
```

```
sentence(S) --> imperative(S), terminator(!).
```

```
/* Noun Phrase */
```

```
np(np(Agmt,Pronoun,[]),Agmt,NPCase,def,_,Set,Nil) -->
```

```
    {is_pp(Set)},
```

```
    pers_pron(Pronoun,Agmt,Case),
```

```
    {empty(Nil), role(Case,decl,NPCase)}.
```

```
/* Prepositional Phrase */
```

```
pp(pp(Prep,Arg),Case,Set,Mask) -->
```

```
    prep(Prep),
```

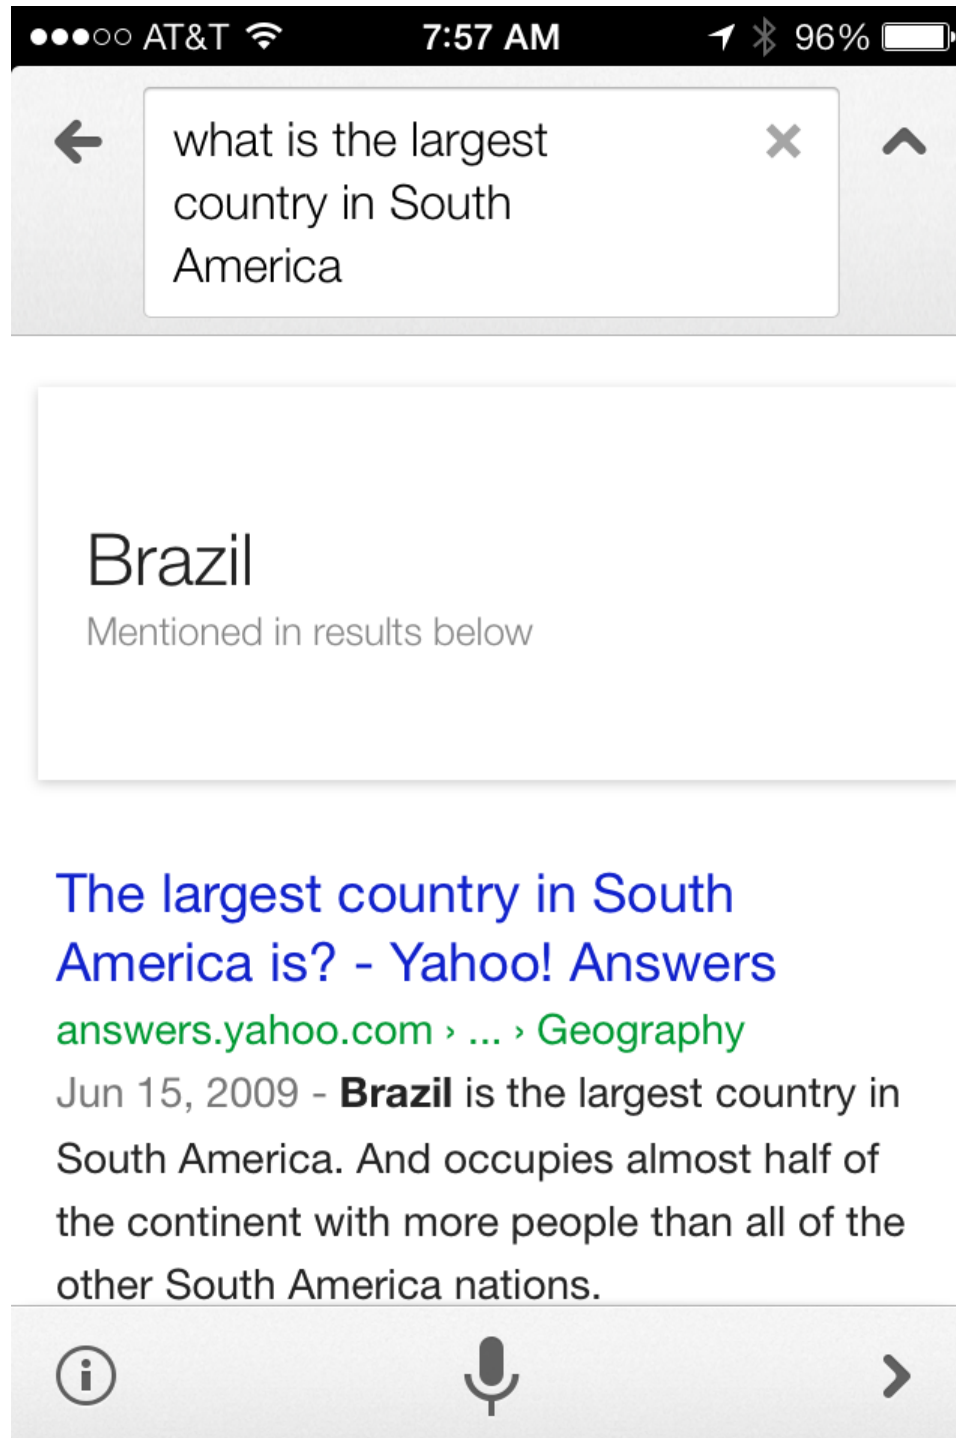
```
    {prep_case(NPCase)},
```

```
    np(Arg,_,NPCase,_,Case,Set,Mask).
```




Fast-forward to 2013

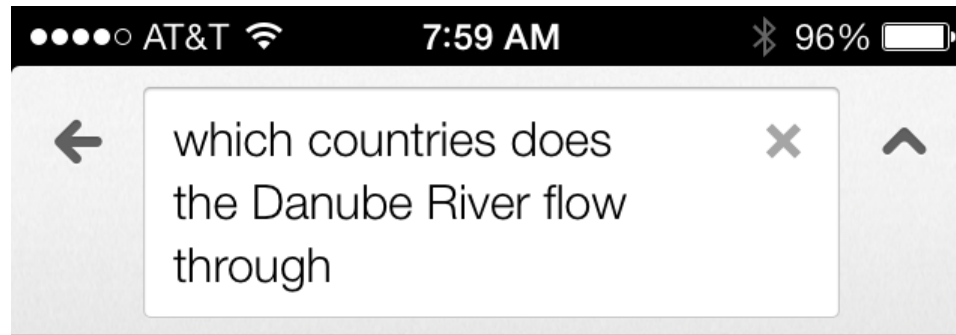
- Siri
- Google now
- Facebook graph search





which countries does
the Danube River flows
through





What Countries Does the Danube River Flow through? - Ask.com


[www.ask.com](#) › ... › Environment

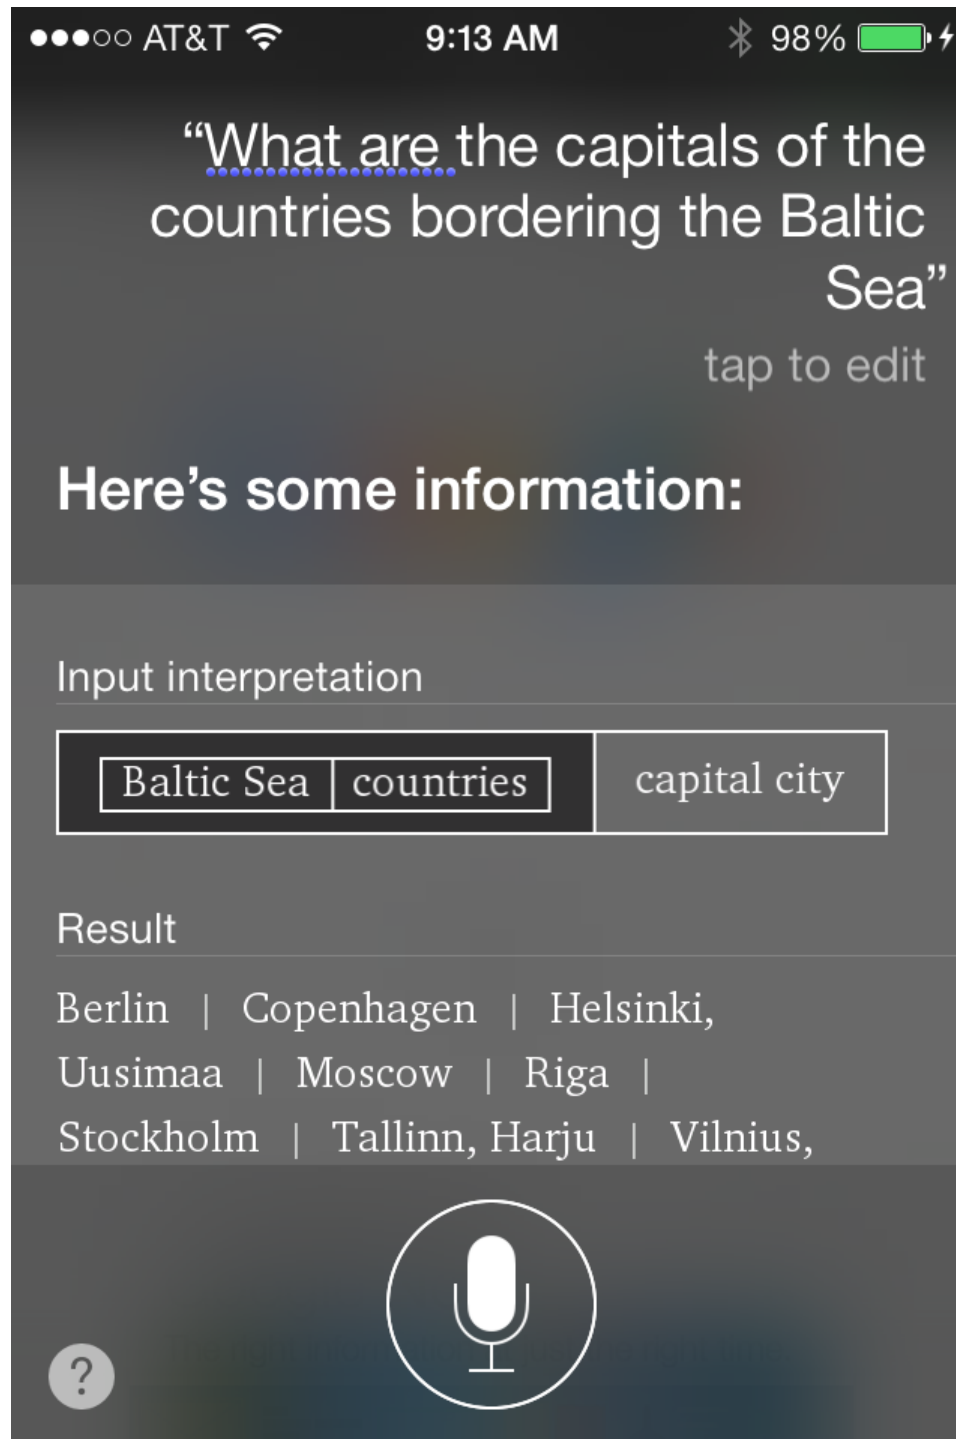
The **Countries** which the **River Danube** flow **through** are; Germany, Austria, Slovakia, Hungary, Croatia, Serbia, Bulgaria ...


What Countries Does the Danube River Run through - Ask.com

[www.ask.com](#) › ... › Environment

River Danube is the second largest in Europe **river** after the **Volga**. It runs in ten **countries** including Germany, Austria, ...









●●●○ AT&T

7:59 AM

96%

I found some tables for four
around 7 pm. Tap the one
you'd like to reserve:

2 Restaurants

Reposado

0.7 mi >

236 Hamilton Ave

Mexican, Vegetarian, \$\$\$\$

Open tables: 6:30 PM, 7:30

★★★★★ 658 Reviews

Joya

0.8 mi >

339 University Ave



Latin American, Mexican, \$\$\$\$






Programming Language Interpreter

- What is meaning of $3+5*6$?
- First parse it into $3+(5*6)$
- Now give a meaning to each node in the tree (bottom-up)






More complex meanings

- How about $3+5*x$?
- Don't know x at compile time
- “Meaning” at a node is a piece of code, not a number


- Form is “rule-to-rule” translation
 - We provide a way to form the semantics of each parent in terms of the semantics of the children





(Formal/Compositional) Computational Semantics

- Sentences: “John smokes.”
“Everyone who smokes snores.”
- Syntactic Analyses:




- Semantics Construction: smoke(j)
 - Logic as meaning representation language
- Inference: $\forall x.\text{smoke}(x) \rightarrow \text{snore}(x)$, smoke(j)
 $\Rightarrow \text{snore}(j)$



Logic: Some Preliminaries

Three major kinds of objects

1. Booleans (Bool)
 - Roughly, the semantic values of sentences
2. Individuals/Entities (Ind)
 - Values of NPs, i.e., objects
 - Maybe also other types of entities, like times
3. Functions of various types
 - A function returning a boolean is called a “predicate”
 - e.g., `frog(x)`, `green(x)`
 - A predicate defines a set of individuals that satisfy it
 - A one argument predicate is called a “property”
 - More complex functions return other functions!
 - Some functions take other functions as arguments!
 - (Higher order functions.)




Logic: Lambda Terms

- Lambda terms:
 - A way of writing “anonymous functions”
 - No function header or function name
 - But defines the key thing: **behavior** of the function
 - Just as we can talk about 3 without naming it “x”
 - Let `square = $\lambda p. p * p$`
 - Equivalent to `int square(p) { return p * p; }`
 - But we can talk about `$\lambda p p * p$` without naming it
 - Format of a lambda term: `λ variable . expression`



Logic: Lambda Terms

- Lambda terms:
 - Let $\text{square} = \lambda p \ p * p$
 - Then $\text{square}(3) = (\lambda p \ p * p)(3) = 3 * 3$
 - **Note:** $\text{square}(x)$ isn't a function! It's just the value $x * x$.
 - But $\lambda x \ \text{square}(x) = \lambda x \ x * x = \lambda p \ p * p = \text{square}$
(proving that these functions are equal – and indeed they are,
as they act the same on all arguments: what is $(\lambda x \ \text{square}(x))(y)?$)
- Let $\text{even} = \lambda p \ (p \bmod 2 == 0)$ a predicate: returns true/false
- $\text{even}(x)$ is true if x is even
- How about $\text{even}(\text{square}(x))$?
- $\lambda x \ \text{even}(\text{square}(x))$ is true of numbers with even squares
 - Just apply rules to get $\lambda x \ (\text{even}(x * x)) = \lambda x \ (x * x \bmod 2 == 0)$
 - This happens to denote the same predicate as even does



Logic: Multiple Arguments

- All lambda terms have one argument
- But we can fake multiple arguments ...
- Suppose we want to write `times(5,6)`
- Remember: `square` can be written as $\lambda x.\text{square}(x)$
- Similarly, `times` is equivalent to $\lambda x.[\lambda y.\text{times}(x,y)]$
- Claim that `times(5)(6)` means same as `times(5,6)`
 - $\text{times}(5) = (\lambda x.\lambda y.\text{times}(x,y)) (5) = \lambda y.\text{times}(5,y)$
 - If this function weren't anonymous, what would we call it?
 - $\text{times}(5)(6) = (\lambda y \text{ times}(5,y))(6) = \text{times}(5,6)$
- Referred to as “**currying**”



Logic: Interesting Constants

- We have “constants” that name some of the entities and functions (e.g., `times`):
 - `GeorgeWBush` – an entity
 - `red` – a predicate on entities
 - holds of just the red entities: `red(x)` is true if `x` is red!
 - `loves` – a predicate on 2 entities
 - `loves(GeorgeWBush, LauraBush)`
 - *Question:* What does `loves(LauraBush)` denote?
- Constants used to define meanings of words
- Meanings of phrases will be built from the constants



Compositional Semantics

- We've discussed what semantic representations should look like.
- **But how do we get them from sentences???**
- **First** - parse to get a syntax tree.
- **Second** - look up the semantics for each word.
- **Third** - build the semantics for each constituent
 - Work from the bottom up
 - The syntax tree is a “recipe” for how to do it
- **Principle of Compositionality**
 - The meaning of a whole is derived from the meanings of the parts, via composition rules



A simple grammar of English

(in Definite Clause Grammar, DCG, form – as in Prolog)

sentence --> noun_phrase, verb_phrase.

noun_phrase --> proper_noun.

noun_phrase --> determiner, noun.

verb_phrase --> verb, noun_phrase.

Proper_noun --> [John]

Proper_noun --> [Mary]

determiner --> [the]

determiner--> [a]

verb --> [ate]

verb --> [kissed]

noun --> [cake]

noun --> [lion]



Extending the grammar to check number agreement between subjects and verbs

$S \rightarrow NP(Num), VP(Num).$

$NP(Num) \rightarrow \text{proper_noun}(Num).$

$NP(Num) \rightarrow \text{det}(Num), \text{noun}(Num).$

$VP(Num) \rightarrow \text{verb}(Num), NP(_).$

$\text{proper_noun}(s) \rightarrow [\text{Mary}].$

$\text{det}(s) \rightarrow [\text{the}].$

$\text{det}(p) \rightarrow [\text{the}].$

$\text{noun}(s) \rightarrow [\text{lion}].$

$\text{noun}(p) \rightarrow [\text{lions}].$

$\text{verb}(s) \rightarrow [\text{eats}].$

$\text{verb}(p) \rightarrow [\text{eat}].$



A simple DCG grammar with semantics

sentence(SMeaning) --> noun_phrase(NPMeaning),
verb_phrase(VPMeaning), {combine (NPMeaning,
VPMeaning, SMeaning)}.

verb_phrase(VPMeaning) --> verb(Vmeaning),
noun_phrase(NPMeaning), {combine (NPMeaning,
VMeaning, VPMeaning)}.

noun_phrase (NPMeaning) --> name(NPMeaning).

name(john) --> [john].

verb(λx .jumps(x)) --> [jumps]


name(mary) --> [mary].


verb(λy . λx .loves(x,y)) -->[loves]

Combine(X, Y, Z) --> apply(Y, X, Z)



Parse tree with associated semantics





In detail: Beta-Reduction

$(\lambda y \lambda x. \text{love}(x, y))[\text{mary}][\text{john}]$

$\beta \Rightarrow (\lambda x. \text{love}(x, \text{mary}))[\text{john}]$

$\beta \Rightarrow \text{love}(\text{john}, \text{mary})$



Formal Compositional Semantics ...

- **Richard Montague**
(1930-1971)
- “... *I reject the contention that an important theoretical difference exists between formal and natural languages ...*”






Augmented CFG Rules

- We can also accomplish this just by attaching semantic formation rules to our syntactic CFG rules

$$A \rightarrow \alpha_1 \dots \alpha_n \quad \{f(\alpha_1.sem, \dots, \alpha_n.sem)\}$$


- This should be read as the semantics we attach to A can be computed from some function applied to the semantics of A's parts.
- The functions/operations permitted in the semantic rules are restricted, falling into two classes
 - Pass the semantics of a daughter up unchanged to the mother
 - Apply (as a function) the semantics of one of the daughters of a node to the semantics of the other daughters



How do things get more complex?

(The former) GRE analytic section

- Six sculptures – C, D, E, F, G, H – are to be exhibited in rooms 1, 2, and 3 of an art gallery.
 - Sculptures C and E may not be exhibited in the same room.
 - Sculptures D and G must be exhibited in the same room.
 - If sculptures E and F are exhibited in the same room, no other sculpture may be exhibited in that room.
 - At least one sculpture must be exhibited in each room, and no more than three sculptures may be exhibited in any room.
- If sculpture D is exhibited in room 3 and sculptures E and F are exhibited in room 1, which of the following may be true?
 1. Sculpture C is exhibited in room 1.
 2. Sculpture H is exhibited in room 1.
 3. Sculpture G is exhibited in room 2.
 4. Sculptures C and H are exhibited in the same room.
 5. Sculptures G and F are exhibited in the same room.



Logic: Interesting Constants

- Generalized Quantifiers
- **most** – a predicate on 2 predicates on entities
 - **most(pig, big)** = “most pigs are big”
 - Equivalently, **most(λx pig(x), λx big(x))**
 - returns true if most of the things satisfying the first predicate also satisfy the second predicate
- similarly for other quantifiers
 - **all(pig,big)** (equivalent to $\forall x$ pig(x) \Rightarrow big(x))
 - **exists(pig,big)** (equivalent to $\exists x$ pig(x) AND big(x))
 - can even build complex quantifiers from English phrases:
 - “between 12 and 75”; “a majority of”; “all but the smallest 2”



Quantifier Order

- Groucho Marx celebrates quantifier order ambiguity:
 - In this country a woman gives birth every 15 min.
Our job is to find that woman and stop her.
 - $\exists \text{woman } (\forall 15\text{min gives-birth-during}(\text{woman}, 15\text{min}))$
 - $\forall 15\text{min } (\exists \text{woman gives-birth-during}(15\text{min}, \text{woman}))$
 - Surprisingly, both are possible in natural language!
 - Which is the joke meaning?
 - (where it's always the same woman)



Scope Needs to be Resolved!

At least one sculpture must be exhibited in each room.

The same sculpture in each room?

No more than three sculptures may be exhibited in any room.

Reading 1: For every room, there are no more than three sculptures exhibited in it.

Reading 2: Only three or less sculptures are exhibited (the rest are not shown).


Reading 3: Only a certain set of three or less sculptures may be exhibited in any room (for the other sculptures there are restrictions in allowable rooms).

- Some readings will be ruled out by being uninformative or by contradicting other statements
- Otherwise we must be content with distributions over scope-resolved semantic forms




An alternative: Semantic Grammars

- A problem with traditional linguistic grammars is that they don't necessarily reflect the semantics in a straightforward way
- You can deal with this by...
 - Fighting with the grammar
 - Complex lambdas and complex terms, etc.
 - Rewriting the grammar to reflect the semantics
 - And in the process give up on some syntactic niceties
 - known as "Semantic grammars"
 - Simple idea, dumb name



Lifer Semantic Grammars

- Example domain—access to DB of US Navy ships
 - S → <present> the <attribute> of <ship>
 - <present> → what is | [can you] tell me
 - <attribute> → length | beam | class
 - <ship> → the <shipname>
 - <shipname> → kennedy | enterprise
 - <ship> → <classname> class ships
 - <classname> → kitty hawk | lafayette
- Example inputs recognized by above grammar:
 - can you tell me the class of the Enterprise*
 - what is the length of Kitty Hawk class ships*
 - Many categories are not "true" syntactic categories
 - Words are recognized by their context rather than category (e.g. *class*)
 - Recognition is strongly directed
 - Strong direction useful for error detection and correction
 - G. Hendrix, E. Sacerdoti, D. Sagalowicz, and J. Slocum. 1978. Developing a natural language interface to complex data. ACM Transactions on Database Systems 3:105-147



Semantic Grammar

- The term semantic grammar refers to the motivation for the grammar rules
 - The technology (plain CFG rules with a set of terminals) is the same as we've been using
 - The good thing about them is that you get exactly the semantic rules you need
 - The bad thing is that you need to develop a new grammar for each new domain
- Typically used in conversational agents in constrained domains
 - Limited vocabulary
 - Limited grammatical complexity
 - Syntactic parsing can often produce all that's needed for semantic interpretation even in the face of "ungrammatical" input – write fragment rules



Semantic Grammars Summary

- Advantages:
 - Efficient recognition of limited domain input
 - Absence of overall grammar allows pattern-matching possibilities for idioms, etc.
 - No separate interpretation phase
 - Strength of top-down constraints allows powerful ellipsis mechanisms

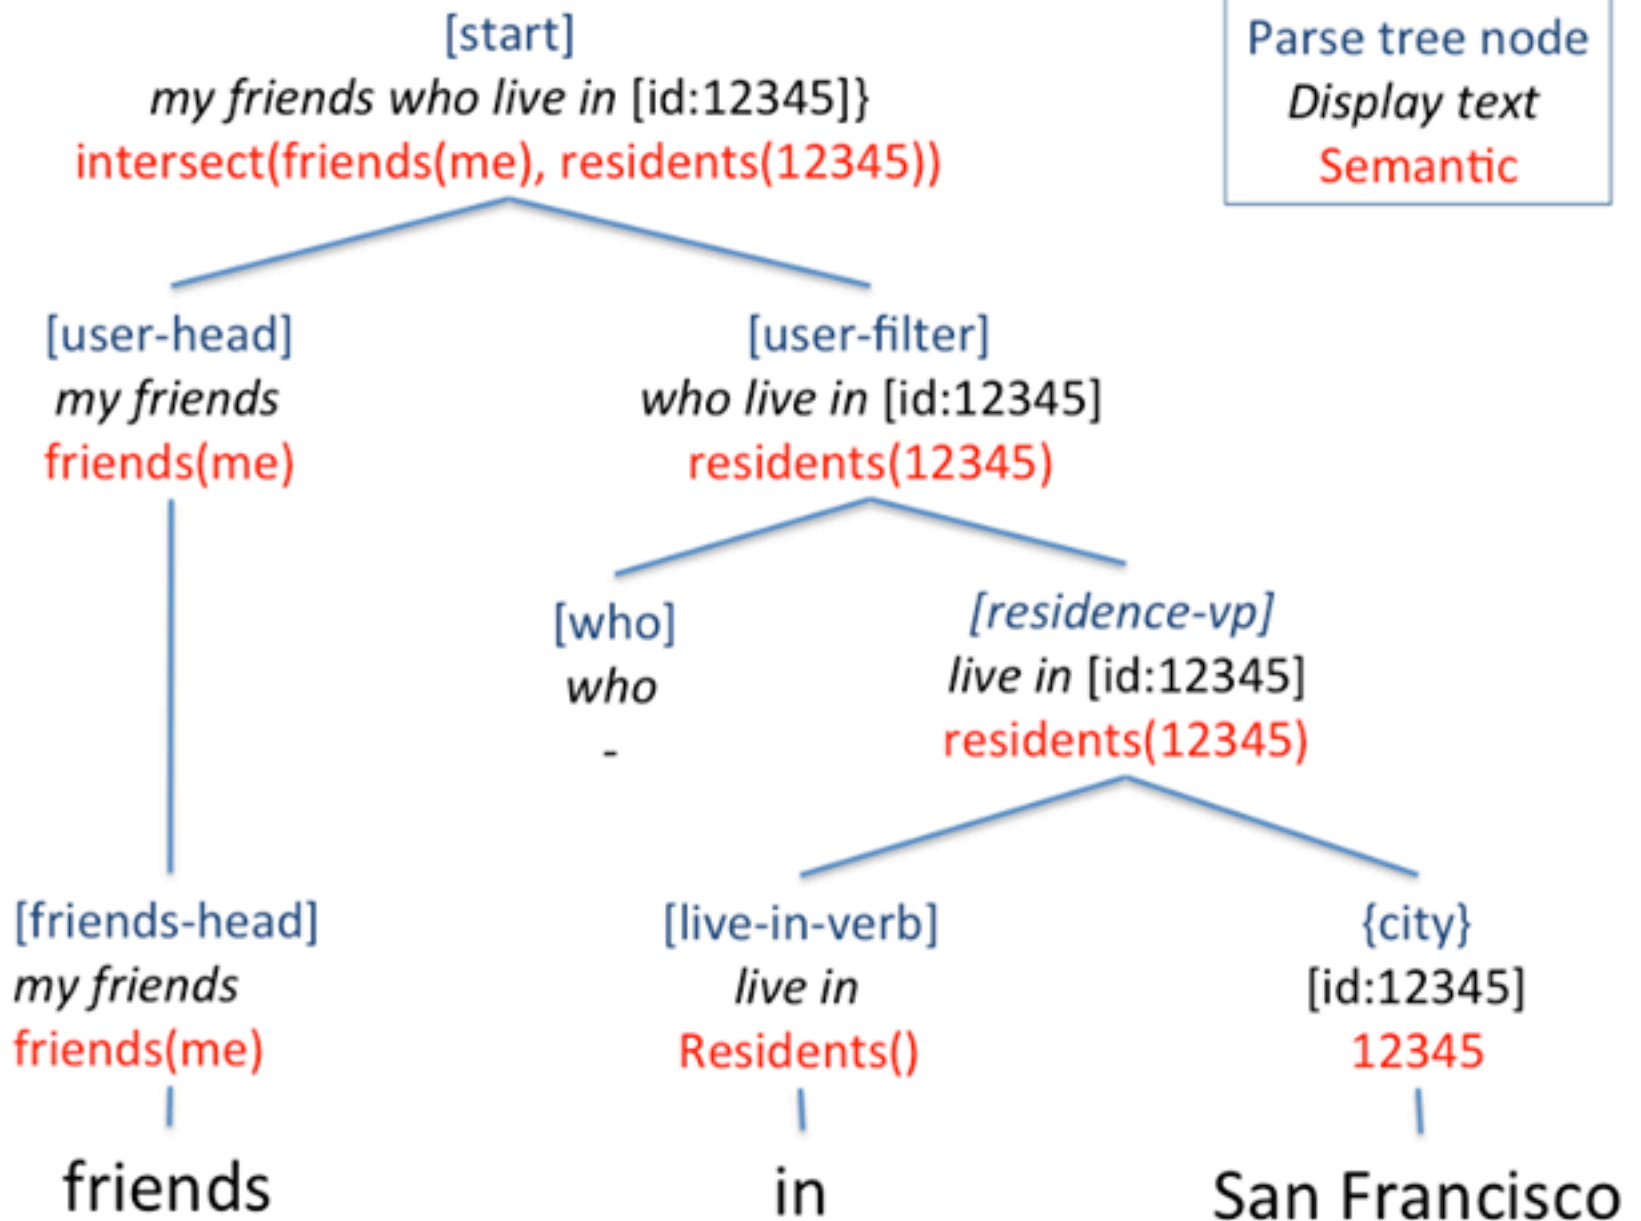
What is the length of the Kennedy? The Kittyhawk?

- Disadvantages:
 - Different grammar required for each new domain
 - Lack of overall syntax can lead to "spotty" grammar coverage
 - E.g. fronting possessive in "<attribute> of <ship>" to <ship> 's
<attribute> doesn't imply fronting in "<rank> of <officer>"
 - Difficult to develop grammars past a certain size
 - Suffers from fragility



Facebook Graph Search

- Uses a weighted context free grammar (WCFG) to represent the Graph Search query language:
- $[start] \Rightarrow [users]$ \$1
- $[users] \Rightarrow \text{my friend}$ friends(me)
- $[users] \Rightarrow \text{friends of } [users]$ friends(\$1)
- $[users] \Rightarrow \{user\}$ \$1
- $[start] \Rightarrow [photos]$ \$1
- $[photos] \Rightarrow \text{photos of } [users]$ photos(\$1)
- A terminal symbol can be an entity, e.g., {user}, {city}, {employer}, {group}; it can also be a word/phrase, e.g., friends, live in, work at, members, etc. A parse tree is produced by starting from [start] and expanding the production rules until it reaches terminal symbols.



The parse tree, semantic and entity ID used in the above example are for illustration only; they do not represent real information used in Graph Search Beta