


MY HOBBY:

SITTING DOWN WITH GRAD STUDENTS AND TIMING
HOW LONG IT TAKES THEM TO FIGURE OUT THAT
I'M NOT ACTUALLY AN EXPERT IN THEIR FIELD.



CS224n is in this interstice

Natural Language Processing

CS224N/Ling284



Christopher Manning
Lecture 1



Course logistics in brief

- Instructor: Christopher Manning
- TAs: Spence Green, Leon Lin, Richard Socher ... and probably more
- Time: MW 11:00–12:15, Skilling Aud ... maybe moving to Gates B03
- Programming language: mainly Java
- Other information: see the class2go webpage
 - <http://cs224n.stanford.edu/>
- “Handouts”: online



This class

- Assumes you come with some skills...
 - Some basic linear algebra, probability, and statistics; decent programming skills
 - But not everyone has the same skills
 - Assumes some ability to learn missing knowledge
- Teaches key theory and methods for (statistical) NLP: MT, information extraction, parsing, semantics, etc.
 - Learn techniques which can be used in practical, robust systems that can (partly) understand human language
- But it's something like an "AI Systems" class:
 - A lot of it is hands-on, problem-based learning
 - Often practical issues are as important as theoretical niceties
 - We often combine a bunch of ideas




Goals of the field of NLP

- Computers would be a lot more useful if they could handle our email, do our library research, chat to us ...
- But they are fazed by natural human languages.
 - Or at least their programmers are ... most people just avoid the problem and get into menus and radio buttons, or XML, or the so-called semantic web, or ...
- But someone has to work on the hard problems!
 - How can we tell computers about language?
 - Or help them to learn it as kids do?



Natural language: the earliest UI

Dave Bowman: Open the pod bay doors, HAL.
HAL: I'm sorry Dave. I'm afraid I can't do that.



(cf. also false Maria in Metropolis – 1926)



What/where is NLP?

- Goals can be very far reaching ...
 - True text understanding and interpretation
 - Real-time participation in spoken dialogs
 - High quality machine translation
- Or very down-to-earth ...
 - Finding the price of products on the web
 - Analyzing reading level or authorship statistically
 - Sentiment detection about products or stocks
 - Extracting names, facts or relations from documents
- These days, the latter predominate
 - As NLP becomes increasingly possible, it becomes increasingly engineering-oriented
 - Also related to changes in approach in AI/NLP in general



Commercial world

the problem session for your course, CS224n in Gates B03 from
1:15-2:05
Skiing Aud

Create New iCal Event...

Show This Date in iCal...



ski resort



Sugar Bowl Ski Lodging

Escape to our Snowbound Village.
Fresh tracks steps from your room.


www.sugarbowl.com

Microsoft®



Their are many approaches,

out



Language Technology

mostly solved

making good progress

still really hard

Spam detection

OK, let's meet by the big ...



D1ck too small? Buy V1AGRA ...



Sentiment analysis

The pho was authentic and yummy.



Waiter ignored us for 20 minutes.



Semantic search

people protesting globalization


Search




...demonstrators stormed IMF offices...

Speech synthesis

I'd like you to listen to me



Coreference resolution



Obama told Mubarak he shouldn't run again.

Question answering (QA)

Q. What currency is used in China?



A. The yuan

Part-of-speech (POS) tagging

ADJ ADJ NOUN VERB ADV

Colorless green ideas sleep furiously.

Word sense disambiguation (WSD)



I need new batteries for my *mouse*.



Textual inference & paraphrase

T. Thirteen soldiers lost their lives ...

H. Several troops were killed in the ...


YES

Named entity recognition (NER)

PERSON ORG LOC

Obama met with UAW leaders in Detroit

Syntactic parsing



I can see Russia from my house!

Summarization

Sheen continues
rant against ...



Sheen is
nuts

Information extraction (IE)

You're invited to our bunga
bunga party, Friday May 27
at 8:30pm in Cordura Hall



Party
May 27
add

Machine translation (MT)

Our specialty is panda fried rice.



我们的专长是熊猫炒饭

Discourse & dialog

Where is Thor playing in SF?



Metreon at 4:30 and 7:30



The hidden structure of language

- We're going beneath the surface...
 - Not just string processing
 - Not just keyword matching in a search engine
 - This is the move that Google has been increasingly engaged in in recent years
 - Moving from matching keywords to satisfying user needs
 - Not just converting a sound stream to a string of words
 - Like Nuance/Google speech recognition
- We want to recover and manipulate at least *some* aspects of language structure and meaning



Is the problem just cycles?

- Bill Gates, Remarks to Gartner Symposium, October 6, 1997:
 - Applications always become more demanding. Until the computer can speak to you in perfect English and understand everything you say to it and learn in the same way that an assistant would learn – until it has the power to do that – we need all the cycles. We need to be optimized to do the best we can. Right now linguistics are right on the edge of what the processor can do. As we get another factor of two, then speech will start to be on the edge of what it can do.





The early history: 1950s

- Earliest
- Foundational
- First
- MT
- Work
- Little
- Series
- Pro





Why NLP is difficult: Newspaper headlines

1. Minister Accused Of Having 8 Wives In Jail
2. Juvenile Court to Try Shooting Defendant
3. Teacher Strikes Idle Kids
4. Miners refuse to work after death
5. Local High School Dropouts Cut in Half
6. Red Tape Holds Up New Bridges
7. Clinton Wins on Budget,
8. Hospitals Are Sued by 7
9. Police: Crack Found in N

Minister Accused Of Having 8 Wives In Jail

May 21, 2007 06:49 AM



ATLANTA (AP) -- A tra
served two years in pri
has been jailed again fo
marry more women.

Bishop Anthony Owens,
Ga., is in a Gwinnett Co
four women claimed he
after being released fro



Why is natural language understanding difficult

Fed raises interest rates 0.5% in effort to control inflation

- *NYT* headline, from better economic times (17 May 2000)



Language: still the ultimate UI



Where is **A Bug's Life** playing in **Mountain View**?

A Bug's Life is playing at the Century 16 Theater.

When is **it** playing **there**?

It's playing at 2pm, 5pm, and 8pm.

OK. I'd like 1 **adult** and 2 **children** for **the first show**.
How much would **that** cost?



But we need **domain knowledge**, **discourse knowledge**, **world knowledge**, linguistic knowledge.



Why is natural language computing hard?

- Natural language is:
 - highly ambiguous at all levels
 - complex and subtle use of context to convey meaning
 - fuzzy, probabilistic
 - involves reasoning about the world
 - a key part of people interacting with other people (a social system):
 - persuading, insulting and amusing them
- But NLP can also be surprisingly easy sometimes:
 - rough text features can often do half the job



OK, why *else* is NLP hard?

Oh so many reasons!

non-standard English

Great job @justinbieber! Were SOO PROUD of what youve accomplished! U taught us 2 #neversaynever & you yourself should never give up either♥

segmentation issues

the New York-New Haven Railroad
the New York-New Haven Railroad

idioms

dark horse
get cold feet
lose face
throw in the towel

neologisms

unfriend
retweet
bromance
teabagger

garden path sentences

The man who hunts ducks out on weekends.
The cotton shirts are made from grows here.

tricky entity names

... a mutation on the *for* gene ...
Where is *A Bug's Life* playing ...
Most of *Let It Be* was recorded ...


world knowledge

Mary and Sue are sisters.
Mary and Sue are mothers.

prosody

I never said *she* stole my money.
I never said she *stole* my money.
I never said she stole *my* money.

lexical specificity



But that's what makes it fun!



Making progress on this problem...

- The task is difficult! What tools do we need?
 - Knowledge about language
 - Knowledge about the world
 - A way to combine knowledge sources
- The answer that's been getting traction:
 - **probabilistic models** built from language data
 - $P(\text{"maison"} \rightarrow \text{"house"})$ **high**
 - $P(\text{"L'avocat général"} \rightarrow \text{"the general avocado"})$ **low**
- Some computer scientists think this is a new "A.I." or "machine learning" idea
 - But really it's an older idea that was taken from the electrical engineers....



Where do we head?

Look at subproblems, approaches, and applications at different levels

- Statistical machine translation
- Statistical NLP: classification and sequence models (part-of-speech tagging, named entity recognition, information extraction)
- Syntactic (probabilistic) parsing
- Building semantic representations from text. QA.
- *(Unfortunately left out: natural language generation, phonology/morphology, speech dialogue systems, more on natural language understanding, There are other classes for some! cs224u/s)*

Machine Translation

美国关岛国际机场及其办公室均接获一名自称沙地阿拉伯富商拉登等发出的电子邮件，威胁将会向机场等公众地方发动生化袭击後，关岛经保持高度戒备。



The U.S. island of Guam is maintaining a high state of alert after the Guam airport and its offices both received an e-mail from someone calling himself the Saudi Arabian Osama bin Laden and threatening a biological/chemical attack against public places such as the airport .

The classic acid test for natural language processing.

Requires capabilities in both interpretation and generation.

About \$26 billion spent annually on human translation.

Scott Klemmer: I learned a surprising fact at our research group lunch today. Google Sketchup releases a version every 18 months, and the primary difficulty of releasing more often is not the difficulty of producing software, but the cost of internationalizing the user manuals!

Many slides from Kevin Knight (at ISI)

Statistical Solution

- Parallel Texts
 - Rosetta Stone

Hieroglyphs

Demotic

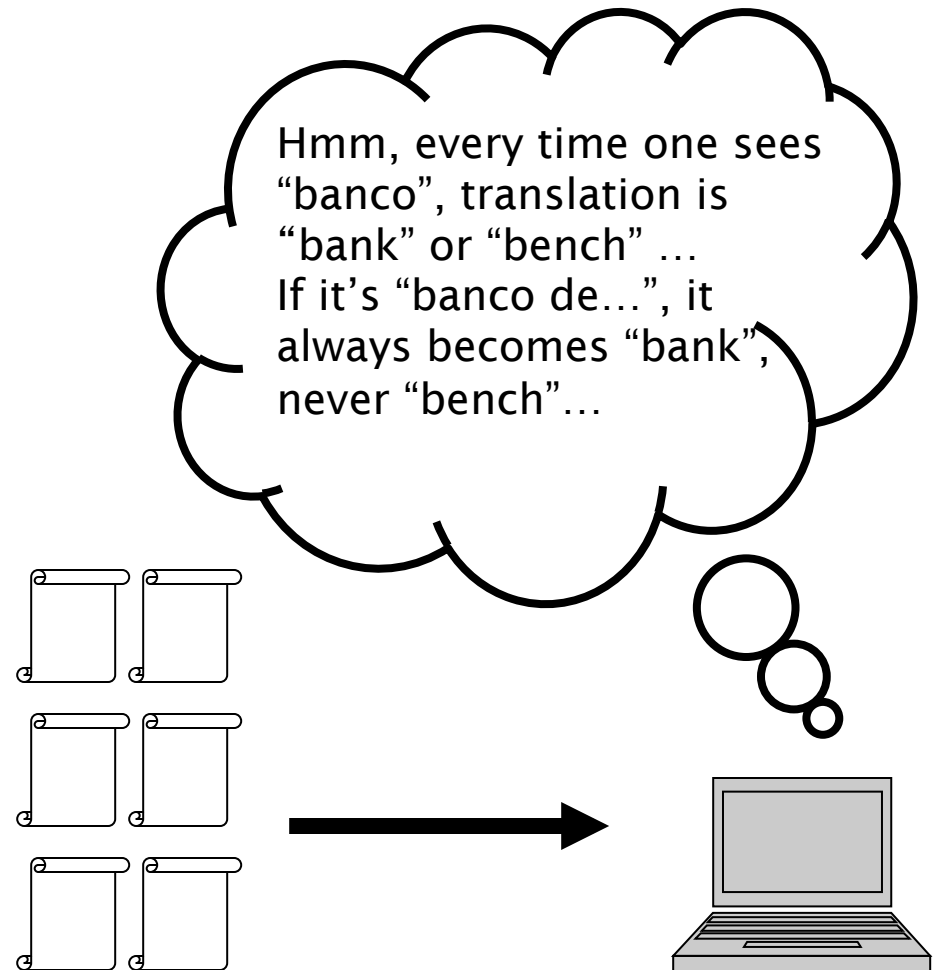
Greek



Statistical Solution

- Parallel Texts

- Instruction Manuals
- Hong Kong/Macao Legislation
- Canadian Parliament Hansards
- United Nations Reports
- Official Journal of the European Communities
- Translated news



Centauri/Arcturan [Knight, 1997]

Your assignment, translate this to Arcturan: farok crrrok hihok yorok klok kantok ok-yurp

Centauri/Arcturan [Knight, 1997]

Your assignment, translate this to Arcturan: farok crrrok hihok yorok klok kantok ok-yurp

1a. ok-voon ororok sprok .	7a. lalok farok ororok lalok sprok izok enemok .
1b. at-voon bichat dat .	7b. wat jjat bichat wat dat vat eneat .
2a. ok-drubel ok-voon anak plok sprok .	8a. lalok brok anak plok nok .
2b. at-drubel at-voon pippat rrat dat .	8b. iat lat pippat rrat nnat .
3a. erok sprok izok hihok ghrok .	9a. wiwok nok izok kantok ok-yurp .
3b. totat dat arrat vat hilat .	9b. totat nnat quat oloat at-yurp .
4a. ok-voon anak drok brok jok .	10a. lalok mok nok yorok ghrok klok .
4b. at-voon krat pippat sat lat .	10b. wat nnat gat mat bat hilat .
5a. wiwok farok izok stok .	11a. lalok nok crrrok hihok yorok zanzanok .
5b. totat jjat quat cat .	11b. wat nnat arrat mat zanzanat .
6a. lalok sprok izok jok stok .	12a. lalok rarok nok izok hihok mok .
6b. wat dat krat quat cat .	12b. wat nnat forat arrat vat gat .